

**CS60002: Distributed Systems**  
**Spring 2019**  
**List of Distributed Systems Projects**

1. Distributed Log Service
  - a. Should support storage and retrieval of record-oriented, append-only logs
  - b. Should tolerate at least one crash failure
  - c. Should be scalable
  - d. Systems to study: Facebook LogDevice, Apache DistributedLog (Part of Apache Bookkeeper now)
2. Distributed Publish Subscribe System
  - a. Should support creation/deletion of new groups
  - b. Should support users to subscribe/unsubscribe to/from one or more groups
  - c. Should support two modes of users – publishers and subscribers. A user can be publisher in some groups and subscribers for others (may be both for a group)
  - d. Should support both push (notification to user) and pull (periodic check by user) modes of events
  - e. Should tolerate at least one crash failure, with eventual consistency for replicas
  - f. Should consider both availability and performance in delivering events to subscribers
  - g. System to study: Apache Kafka
3. Distributed Data Store
  - a. Users should be able to upload/download/search key-value pairs in the store
  - b. System should be able to tolerate one crash fault
  - c. Should ensure sequential consistency if replication is used
  - d. Should be able to add/delete servers from the system
  - e. Should do load balancing
  - f. System to study: Apache Cassandra
4. Distributed Coordination Middleware
  - a. Should support setting up a cluster of nodes for an application
  - b. Should support addition/deletion of nodes
  - c. Should support naming of nodes
  - d. Should support locking of shared data
  - e. Should support atomic actions on shared/replicated data even in the presence of on crash fault
  - f. Project should also develop any simple application of its choice over the middleware using the middleware APIs
  - g. System to study: Apache Zookeeper
5. Distributed Directory Service

- a. Should support definition/maintenance/updation of object schema (what classes of objects with what attributes).
  - b. Should support both must-have and may-have attributes for objects
  - c. Should support addition/deletion/modification of objects with schema check
  - d. Should support base, one-level, and subtree search for LDAP
  - e. Should support the directory tree to be partitioned into at least 2 parts and stored separately
  - f. Should be tolerant to at least one crash fault (with sequential consistency for schema change and eventual consistency for object change)
  - g. Should support consistent merge with timestamps
  - h. Systems to study: Apache Directory, OpenLDAP
- 6. Cluster Middleware with Central Master Node
  - a. Should support addition/deletion of compute nodes (assume homogeneous)
  - b. Should support job submission/withdrawal/status-check at any time
  - c. Should support load balancing of jobs across compute nodes
  - d. Should be tolerant to at least one crash fault (for both master and compute nodes)
  - e. Should support job policies in master node
  - f. Systems to study: PBS Pro, ROCKS, Google Borg
- 7. Cluster Middleware with No Central Master Node
  - a. Should support addition/deletion of compute nodes of any types (heterogeneous)
  - b. Should support job submission/withdrawal/status-check at any time at any node
  - c. Should support discovery of compute nodes
  - d. Should support matchmaking of jobs with compute nodes
  - e. Should be tolerant to at least one crash fault
  - f. Systems to study: HTCondor, Globus
- 8. Multimedia Content Delivery Network
  - a. Should support upload/download/search of multimedia files
  - b. Should support delivery of multimedia contents to users spread geographically
  - c. Should support both static replicas for fault tolerance (one crash fault) and dynamic replicas based on demand for performance.
  - d. Paper: to study:
  - e. System to study: Akamai
- 9. HTTP Content Delivery
  - a. Should support faster download of webpages through the services of a CDN
  - b. Should support tolerance to at least one crash fault in the CDN
  - c. Should try to optimize network use
  - d. Should support load balancing between web servers
  - e. System to study: Cloudflare

#### 10. Distributed File System

- a. Should be able to mount remote filesystems locally to give an integrated view locally
- b. Should be able to read and write both local and remote files with the same interface (use FUSE)
- c. Should support locking of shared files
- d. Should support caching for performance (define consistency to be supported)
- e. Should tolerate at least one crash fault
- f. Systems to study: NFS

#### 11. Distributed chat

- a. Users should be able to send and receive messages
- b. All users should see the same sequence of messages (total order)
- c. Causal relationship of messages to be maintained (causal order)
- d. No centralized server
- e. Users should be able to join and leave the chat anytime (new users should see all messages that originate after the user join)
- f. Systems to study: Orbit Web, Tox

#### 12. Collaborative Download

- a. Should support discovery of content already present nearby (can assume a closed group)
- b. Should support entry/exit of users into the group
- c. Should support download partitioning based on load balancing/incentives
- d. Reading Material: [paper-1](#) [paper-2](#)

#### 13. P2P File Sharing System

- a. Users should be able to join and leave system asynchronously.
- b. Should be scalable with no single point of failure.
- c. Users should be able to search, upload and download files.
- d. Must ensure successful file download if file present in the system.
- e. Systems to study: dc++, Gnutella, bit torrent

#### 14. Software Load Balancer

- a. Should distribute the workload between the active servers fairly.
- b. Should reduce single point of failure through redundancy and rebalance workloads when a server fails.
- c. Improve scalability by the addition of new servers dynamically.
- d. Should be resilient to single crash fault on itself.
- e. Systems to study: Seesaw, HAProxy

#### 15. Railway Reservation System

- a. Should consider type and dynamicity of data for data replication with different consistency models
- b. Should tolerate at least one crash failure

- c. Should support load balancing between backend servers

## 16. Travel Agency Booking System