

Maekawa's Algorithm

- Permission obtained from only a subset of other processes, called the *Request Set* (or *Quorum*)
- Separate Request Set, R_i , for each process i
- Requirements:
 - for all i, j : $R_i \cap R_j \neq \Phi$
 - for all i : $i \in R_i$
 - for all i : $|R_i| = K$, for some K
 - any node i is contained in exactly D Request Sets, for some D
- **$K = D = \sqrt{N}$** for Maekawa's

A Simple Version

- To request critical section:
 - i sends REQUEST message to all process in R_i
- On receiving a REQUEST message:
 - Send a REPLY message if no REPLY message has been sent since the last RELEASE message is received.
 - Update status to indicate that a REPLY has been sent.
 - Otherwise, queue up the REQUEST
- To enter critical section:
 - i enters critical section after receiving REPLY from all nodes in R_i

A Simple Version contd..

- To release critical section:
 - Send RELEASE message to all nodes in R_i
 - On receiving a RELEASE message, send REPLY to next node in queue and delete the node from the queue.
 - If queue is empty, update status to indicate no REPLY message has been sent.

Features

- Message Complexity: $3 * \sqrt{N}$
- Synchronization delay =
 - $2 * (\text{max message transmission time})$
- Major problem: DEADLOCK possible
- Need three more types of messages (FAILED, INQUIRE, YIELD) to handle deadlock.
 - Message complexity can be $5 * \text{sqrt}(N)$
- Building the request sets?

Token based Algorithms

- Single token circulates, enter CS when token is present
- Mutual exclusion obvious
- Algorithms differ in how to find and get the token
- Uses sequence numbers rather than timestamps to differentiate between old and current requests

Suzuki Kasami Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it
- Issues:
 - Current versus outdated requests
 - Determining sites with pending requests
 - Deciding which site to give the token to

Suzuki Kasami Algorithm

- The token:
 - Queue (FIFO) Q of requesting processes
 - $LN[1..n]$: sequence number of request that j executed most recently
- The request message:
 - $REQUEST(i, k)$: request message from node i for its k^{th} critical section execution
- Other data structures
 - $RN_i[1..n]$ for each node i , where $RN_i[j]$ is the largest sequence number received so far by i in a REQUEST message from j .

Suzuki Kasami Algorithm

- To request critical section:
 - If i does not have token, increment $RN_i[i]$ and send $REQUEST(i, RN_i[i])$ to all nodes
 - If i has token already, enter critical section if the token is idle (no pending requests), else follow rule to release critical section
- On receiving $REQUEST(i, sn)$ at j :
 - Set $RN_j[i] = \max(RN_j[i], sn)$
 - If j has the token and the token is idle, then send it to i if $RN_j[i] = LN[i] + 1$. If token is not idle, follow rule to release critical section

Suzuki Kasami Algorithm

- To enter critical section:
 - Enter CS if token is present
- To release critical section:
 - Set $LN[i] = RN_i[i]$
 - For every node j which is not in Q (in token), add node j to Q if $RN_i[j] = LN[j] + 1$
 - If Q is non empty after the above, delete first node from Q and send the token to that node

Notable features

- No. of messages:
 - 0 if node holds the token already, n otherwise
- Synchronization delay:
 - 0 (node has the token) or max. message delay (token is elsewhere)
- No starvation

Raymond's Algorithm

- Forms a directed tree (logical) with the token-holder as root
- Each node has variable "*Holder*" that points to its parent on the path to the root.
 - Root's Holder variable points to itself
- Each node i has a FIFO request queue Q_i