

# Maekawa's Algorithm

- Permission obtained from only a subset of other processes, called the *Request Set* (or *Quorum*)
- Separate Request Set,  $R_i$ , for each process  $i$
- Requirements:
  - for all  $i, j$ :  $R_i \cap R_j \neq \Phi$
  - for all  $i$ :  $i \in R_i$
  - for all  $i$ :  $|R_i| = K$ , for some  $K$
  - any node  $i$  is contained in exactly  $D$  Request Sets, for some  $D$
- **$K = D = \sqrt{N}$**  for Maekawa's

# A Simple Version

- To request critical section:
  - $i$  sends REQUEST message to all process in  $R_i$
- On receiving a REQUEST message:
  - Send a REPLY message if no REPLY message has been sent since the last RELEASE message is received.
  - Update status to indicate that a REPLY has been sent.
  - Otherwise, queue up the REQUEST
- To enter critical section:
  - $i$  enters critical section after receiving REPLY from all nodes in  $R_i$

# A Simple Version contd..

- To release critical section:
  - Send RELEASE message to all nodes in  $R_i$
  - On receiving a RELEASE message, send REPLY to next node in queue and delete the node from the queue.
  - If queue is empty, update status to indicate no REPLY message has been sent.

# Features

- Message Complexity:  $3 * \sqrt{N}$
- Synchronization delay =
  - $2 * (\text{max message transmission time})$
- Major problem: DEADLOCK possible
- Need three more types of messages (FAILED, INQUIRE, YIELD) to handle deadlock.
  - Message complexity can be  $5 * \text{sqrt}(N)$
- Building the request sets?

# Token based Algorithms

- Single token circulates, enter CS when token is present
- Mutual exclusion obvious
- Algorithms differ in how to find and get the token
- Uses sequence numbers rather than timestamps to differentiate between old and current requests

# Suzuki Kasami Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it
- Issues:
  - Current versus outdated requests
  - Determining sites with pending requests
  - Deciding which site to give the token to

# Suzuki Kasami Algorithm

- The token:
  - Queue (FIFO)  $Q$  of requesting processes
  - $LN[1..n]$  : sequence number of request that  $j$  executed most recently
- The request message:
  - $REQUEST(i, k)$ : request message from node  $i$  for its  $k^{\text{th}}$  critical section execution
- Other data structures
  - $RN_i[1..n]$  for each node  $i$ , where  $RN_i[j]$  is the largest sequence number received so far by  $i$  in a REQUEST message from  $j$ .

# Suzuki Kasami Algorithm

- To request critical section:
  - If  $i$  does not have token, increment  $RN_i[i]$  and send  $REQUEST(i, RN_i[i])$  to all nodes
  - If  $i$  has token already, enter critical section if the token is idle (no pending requests), else follow rule to release critical section
- On receiving  $REQUEST(i, sn)$  at  $j$ :
  - Set  $RN_j[i] = \max(RN_j[i], sn)$
  - If  $j$  has the token and the token is idle, then send it to  $i$  if  $RN_j[i] = LN[i] + 1$ . If token is not idle, follow rule to release critical section



# Suzuki Kasami Algorithm

- To enter critical section:
  - Enter CS if token is present
- To release critical section:
  - Set  $LN[i] = RN_i[i]$
  - For every node  $j$  which is not in  $Q$  (in token), add node  $j$  to  $Q$  if  $RN_i[j] = LN[j] + 1$
  - If  $Q$  is non empty after the above, delete first node from  $Q$  and send the token to that node

# Notable features

- No. of messages:
  - 0 if node holds the token already,  $n$  otherwise
- Synchronization delay:
  - 0 (node has the token) or max. message delay (token is elsewhere)
- No starvation

# Raymond's Algorithm

- Forms a directed tree (logical) with the token-holder as root
- Each node has variable "*Holder*" that points to its parent on the path to the root.
  - Root's Holder variable points to itself
- Each node  $i$  has a FIFO request queue  $Q_i$

# Raymond's Algorithm

- To request critical section:
  - Send REQUEST to parent on the tree, provided  $i$  does not hold the token currently and  $Q_i$  is empty. Then place request in  $Q_i$
- When a non-root node  $j$  receives a request from  $i$ 
  - place request in  $Q_j$
  - send REQUEST to parent if no previous REQUEST sent

# Raymond's Algorithm

- When the root receives a REQUEST:
  - send the token to the requesting node
  - set *Holder* variable to point to that node
- When a node receives the token:
  - delete first entry from the queue
  - send token to that node
  - set *Holder* variable to point to that node
  - if queue is non-empty, send a REQUEST message to the parent (node pointed at by *Holder* variable)

# Raymond's Algorithm

- To execute critical section:
  - enter if token is received and own entry is at the top of the queue;  
delete the entry from the queue
- To release critical section
  - if queue is non-empty, delete first entry from the queue, send token to that node and make *Holder* variable point to that node
  - If queue is still non-empty, send a REQUEST message to the parent (node pointed at by *Holder* variable)

# Notable features

- Average message complexity:  $O(\log n)$
- Sync. delay =  $(T \log n)/2$ , where  $T$  = max. message delay